

# GARNET


## Final Presentation

Casper Pranger, Marie Bocher  
William Sawyer, Andreas Herten

# About GARNET

- Is built on top of PETSc, making use of its PC, KSP, and SNES objects.
- Provides automatic space/time discretization with staggered grid FD (1D/2D/3D).
- **Restriction:** (logically) cartesian rectangular.
- **Lightweight:** communication pattern only ever unidirectional 1D. Arithmetic indexing.
- Provides a toolbox of (tensor-valued) fields and differential operators.
- Implements only numerics, leaves physics to the user.
  - **Goal:** enabling fast development and execution of physics and algorithms

GARNET



# Employed Technologies

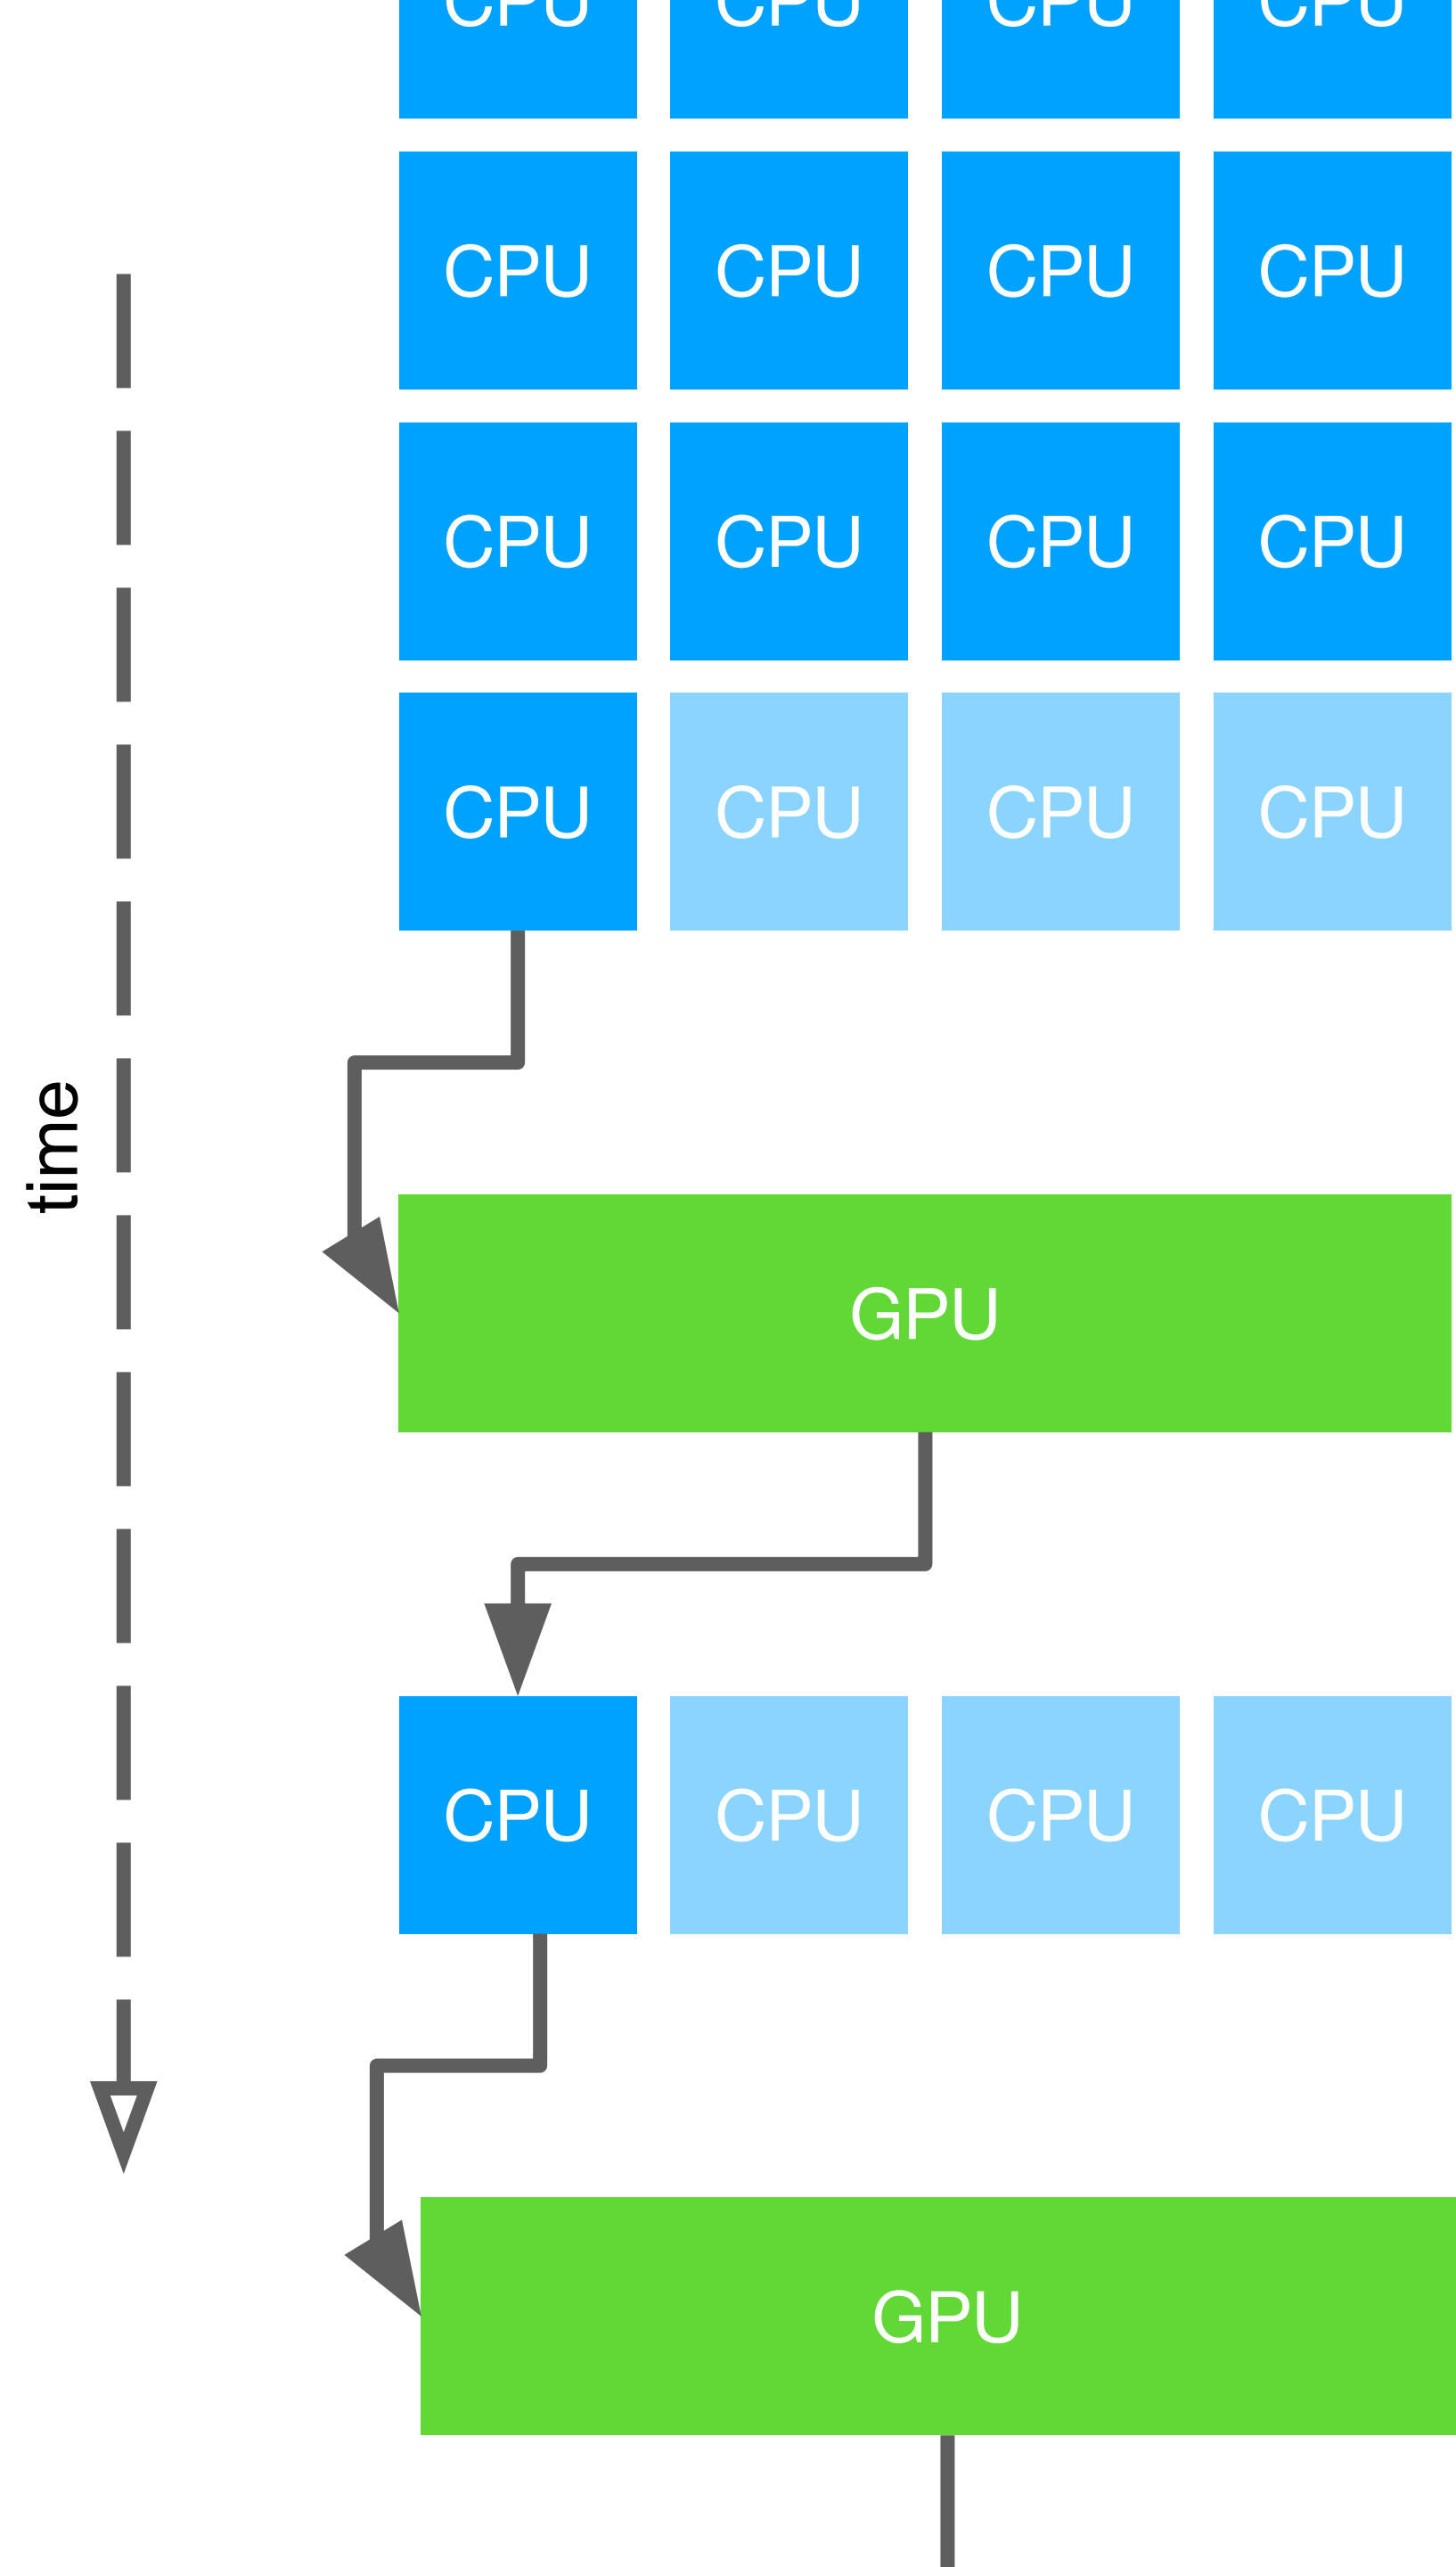
- **CPU:**

- C++14 (Variadic templates, lambdas, iterators, ...)
- MPI 3 Shared Memory & P2P

- **GPU:** same as above (also variadic templates!), plus **Thrust**

- CUDA 9.0
- GCC 6.2.0 (latest version supported by nvcc)






- Multiple ranks (with MPI-shared memory) run on CPU data
- Before GPU invocation: MPI Barrier
- Rank 0 launches GPU computation on full memory
- After GPU computation conclusion: MPI Barrier
- Continue run of multiple MPI ranks on CPU data

## Source Code Snippet – user interface


```
//////////  
//          //  
//  PHYSICS  //  
//          //  
//////////  
  
// const double dt,  
auto bulk_constitutive =  
    [=] __host__ __device__ ( double v_1, double DeltaP )  
    { return v_1 - dt * DeltaP / rho; };  
  
// const double dt,  
auto momentum_balance =  
    [=] __host__ __device__ ( double P_1, double Deltav )  
    { return P_1 - dt * Deltav / beta; };
```

GARNET



## Source Code Snippet – user interface

```
////////////////////////////////////  
//                               //  
//  INITIAL CONDITIONS          //  
//                               //  
////////////////////////////////////  
  
auto gaussian_pulse =  
    [=] __host__ __device__ ( gpu::array<double,N> x )  
    { return std::exp(-0.5*(x[0]*x[0]+x[1]*x[1])/(dx*dx)); };  
  
P_1.SetCoordinated( gaussian_pulse );
```



## Source Code Snippet – backend


```
template< size_t N >
template< class F, class T, size_t ... I >
void Node<N>::set_impl( F&& f, T args, std::index_sequence<I...> )
{
    ...

    auto assignment_f = [=] __host__ __device__ ( arg_tuple_t args ) -> bool
    { thrust::get<0>(args) = f( thrust::get<I+1>(args)... ); return true; };

    grid->amorphous_loop( domain, assignment_f, *this,
std::forward<decltype(std::get<I>(args))>(std::get<I>(args))... );

    ...
}
```

GARNET



# Source Code Snippet – backend


```
template< size_t _N >
template< class F, class ... A >
void Grid<_N>::amorphous_loop( Domain<_N>* domain, F&& f, A&&... args ) const
{
    auto arg_begin_iterator = thrust::make_zip_iterator(
        thrust::make_tuple(
            thrust_iteratoror<std::decay_t<A>,N>(args).d_begin()...
        )
    );

    auto arg_end_iterator    = thrust::make_zip_iterator(
        thrust::make_tuple(
            thrust_iteratoror<std::decay_t<A>,N>(args).d_end()...
        )
    );

    auto discard = thrust::make_discard_iterator();


    MPI_Barrier( domain->get_shared_comm() );
    if( mpi::get_rank( domain->get_shared_comm() ) == 0 )
        thrust::transform( arg_begin_iterator, arg_end_iterator, discard, f );
    MPI_Barrier( domain->get_shared_comm() );
}
```


GARNET





# Speedup

- Currently: 40×
  - (1 GPU) vs. (CPU with 12 MPI ranks (1 node))
  - Simple example which exclusively runs on GPU
  - *But something strange is going on in CPU code* 




GARNET



# Future Work

- GPU-capable boundary conditions, MPI P2P halo exchange
- Compose stencil iterators with point-wise iterators at compile-time
- Move back up to real applications
  - complicated physics => more computational intensity on GPU
- Merge CPU- and GPU-capable versions into one code -- How?

GARNET



# Comments / Feedback

- Many issues of C++14 compatibility with compilers, debuggers, mentors, libraries...

Will improve with time? But sooner rather than later ;-)

- Thrust on Github is not the latest version of Thrust
  - Make Thrust variadic already!
- Error messages are cryptic!
  - Let Thrust produce better error message (is that even possible?)
  - Offer CUDA-compatible Clang on Piz Daint
- Hackathon could be maybe two weeks long :-)
- We'll be back! -- would be far more productive with gpu working from the outset

GARNET

